

```
pip install xgboost imbalanced-learn

Defaulting to user installation because normal site-packages is not
writeable
Collecting xgboost
  Downloading xgboost-3.0.0-py3-none-win_amd64.whl.metadata (2.1 kB)
Requirement already satisfied: imbalanced-learn in c:\programdata\
anaconda3\lib\site-packages (0.12.3)
Requirement already satisfied: numpy in c:\programdata\anaconda3\lib\
site-packages (from xgboost) (1.26.4)
Requirement already satisfied: scipy in c:\programdata\anaconda3\lib\
site-packages (from xgboost) (1.13.1)
Requirement already satisfied: scikit-learn>=1.0.2 in c:\programdata\
anaconda3\lib\site-packages (from imbalanced-learn) (1.5.1)
Requirement already satisfied: joblib>=1.1.1 in c:\programdata\
anaconda3\lib\site-packages (from imbalanced-learn) (1.4.2)
Requirement already satisfied: threadpoolctl>=2.0.0 in c:\programdata\
anaconda3\lib\site-packages (from imbalanced-learn) (3.5.0)
  Downloading xgboost-3.0.0-py3-none-win_amd64.whl (150.0 MB)
----- 0.0/150.0 MB ? eta ------
----- 0.3/150.0 MB ? eta ------
----- 3.9/150.0 MB 13.1 MB/s eta
0:00:12
----- 8.1/150.0 MB 15.7 MB/s eta
0:00:10
----- 12.6/150.0 MB 17.1 MB/s
eta 0:00:09
----- 17.6/150.0 MB 18.8 MB/s
eta 0:00:08
----- 21.2/150.0 MB 18.6 MB/s
eta 0:00:07
----- 26.7/150.0 MB 19.7 MB/s
eta 0:00:07
----- 30.9/150.0 MB 19.6 MB/s
eta 0:00:07
----- 35.1/150.0 MB 19.7 MB/s
eta 0:00:06
----- 38.5/150.0 MB 19.3 MB/s
eta 0:00:06
----- 41.4/150.0 MB 19.5 MB/s
eta 0:00:06
----- 46.9/150.0 MB 19.4 MB/s
eta 0:00:06
----- 51.6/150.0 MB 19.7 MB/s
eta 0:00:05
----- 56.9/150.0 MB 20.0 MB/s
eta 0:00:05
----- 62.1/150.0 MB 20.3 MB/s
eta 0:00:05
----- 67.1/150.0 MB 20.6 MB/s
```

```
eta 0:00:05          ----- 71.3/150.0 MB 20.5 MB/s
eta 0:00:04          ----- 75.8/150.0 MB 20.4 MB/s
eta 0:00:04          ----- 80.2/150.0 MB 20.5 MB/s
eta 0:00:04          ----- 84.4/150.0 MB 20.5 MB/s
eta 0:00:04          ----- 88.3/150.0 MB 20.4 MB/s
eta 0:00:04          ----- 92.3/150.0 MB 20.2 MB/s
eta 0:00:03          ----- 95.9/150.0 MB 20.1 MB/s
eta 0:00:03          ----- 99.1/150.0 MB 20.1 MB/s
eta 0:00:03          ----- 102.8/150.0 MB 19.7 MB/s
eta 0:00:03          ----- 105.9/150.0 MB 19.5 MB/s
eta 0:00:03          ----- 109.3/150.0 MB 19.3 MB/s
eta 0:00:03          ----- 112.5/150.0 MB 19.2 MB/s
eta 0:00:02          ----- 115.9/150.0 MB 19.1 MB/s
eta 0:00:02          ----- 119.0/150.0 MB 18.9 MB/s
eta 0:00:02          ----- 122.4/150.0 MB 18.8 MB/s
eta 0:00:02          ----- 125.3/150.0 MB 18.7 MB/s
eta 0:00:02          ----- 128.7/150.0 MB 18.6 MB/s
eta 0:00:02          ----- 132.1/150.0 MB 18.5 MB/s
eta 0:00:01          ----- 135.5/150.0 MB 18.4 MB/s
eta 0:00:01          ----- 139.2/150.0 MB 18.3 MB/s
eta 0:00:01          ----- 142.6/150.0 MB 18.3 MB/s
eta 0:00:01          ----- 146.0/150.0 MB 18.2 MB/s
eta 0:00:01          ----- 149.7/150.0 MB 18.2 MB/s
eta 0:00:01          ----- 149.9/150.0 MB 18.2 MB/s
eta 0:00:01
```

```

----- 150.0/150.0 MB 17.5 MB/s
eta 0:00:00
Installing collected packages: xgboost
Successfully installed xgboost-3.0.0
Note: you may need to restart the kernel to use updated packages.

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px
import plotly.graph_objects as go
from plotly.subplots import make_subplots
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report, confusion_matrix,
accuracy_score, precision_recall_curve, roc_curve, auc
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from xgboost import XGBClassifier
from imblearn.over_sampling import SMOTE

import warnings
warnings.filterwarnings('ignore')

plt.style.use('fivethirtyeight')
sns.set_palette('viridis')

df = pd.read_csv(r"C:\EXCEL_FILES\PS_20174392719_1491204439457_log.csv")
print(f"Dataset shape: {df.shape}")
df.head()

Dataset shape: (6362620, 11)

   step      type    amount   nameOrig  oldbalanceOrg
newbalanceOrig \
0       1  PAYMENT  9839.64  C1231006815        170136.0
160296.36
1       1  PAYMENT  1864.28  C1666544295        21249.0
19384.72
2       1  TRANSFER   181.00  C1305486145         181.0
0.00
3       1  CASH_OUT   181.00  C840083671         181.0
0.00
4       1  PAYMENT  11668.14  C2048537720        41554.0
29885.86

```

```

      nameDest  oldbalanceDest  newbalanceDest  isFraud
isFlaggedFraud
0   M1979787155           0.0           0.0       0
0
1   M2044282225           0.0           0.0       0
0
2   C553264065            0.0           0.0       1
0
3   C38997010            21182.0          0.0       1
0
4   M1230701703           0.0           0.0       0
0

print("Dataset Information:")
df.info()

print("\nSummary Statistics:")
print(df.describe().T)

missing_values = df.isnull().sum()
print("\nMissing values in each column:")
print(missing_values[missing_values > 0] if missing_values.sum() > 0
else "No missing values")

Dataset Information:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6362620 entries, 0 to 6362619
Data columns (total 11 columns):
 #   Column            Dtype  
--- 
 0   step              int64  
 1   type              object  
 2   amount             float64
 3   nameOrig          object  
 4   oldbalanceOrg     float64
 5   newbalanceOrig    float64
 6   nameDest           object  
 7   oldbalanceDest    float64
 8   newbalanceDest    float64
 9   isFraud            int64  
 10  isFlaggedFraud   int64  
dtypes: float64(5), int64(3), object(3)
memory usage: 534.0+ MB

Summary Statistics:
                count        mean         std      min
25% \step      6362620.0  2.433972e+02  1.423320e+02  1.0  156.00
amount        6362620.0  1.798619e+05  6.038582e+05  0.0  13389.57

```

oldbalanceOrg	6362620.0	8.338831e+05	2.888243e+06	0.0	0.00
newbalanceOrig	6362620.0	8.551137e+05	2.924049e+06	0.0	0.00
oldbalanceDest	6362620.0	1.100702e+06	3.399180e+06	0.0	0.00
newbalanceDest	6362620.0	1.224996e+06	3.674129e+06	0.0	0.00
isFraud	6362620.0	1.290820e-03	3.590480e-02	0.0	0.00
isFlaggedFraud	6362620.0	2.514687e-06	1.585775e-03	0.0	0.00

	50%	75%	max
step	239.000	3.350000e+02	7.430000e+02
amount	74871.940	2.087215e+05	9.244552e+07
oldbalanceOrg	14208.000	1.073152e+05	5.958504e+07
newbalanceOrig	0.000	1.442584e+05	4.958504e+07
oldbalanceDest	132705.665	9.430367e+05	3.560159e+08
newbalanceDest	214661.440	1.111909e+06	3.561793e+08
isFraud	0.000	0.000000e+00	1.000000e+00
isFlaggedFraud	0.000	0.000000e+00	1.000000e+00

Missing values in each column:

No missing values

```

fraud_distribution = df['isFraud'].value_counts(normalize=True) * 100
print("Distribution of transactions:")
print(fraud_distribution)

plt.figure(figsize=(10, 6))
sns.countplot(x='isFraud', data=df)
plt.title('Distribution of Fraudulent vs Non-Fraudulent Transactions')
plt.xticks([0, 1], ['Non-Fraudulent', 'Fraudulent'])
plt.ylabel('Count')
plt.show()


```

Distribution of transactions:

```

isFraud
0    99.870918
1     0.129082
Name: proportion, dtype: float64

```



```

transaction_types = df['type'].value_counts()
print("Transaction Types:")
print(transaction_types)


plt.figure(figsize=(12, 6))
sns.countplot(x='type', data=df, hue='isFraud')
plt.title('Transaction Types and Fraud Distribution')
plt.xlabel('Transaction Type')
plt.ylabel('Count')
plt.xticks(rotation=45)
plt.legend(['Legitimate', 'Fraudulent'])
plt.show()

fraud_by_type = df.groupby('type')['isFraud'].mean() * 100
print("Fraud Rate by Transaction Type (%):\n", fraud_by_type)


plt.figure(figsize=(12, 6))
fraud_by_type.sort_values(ascending=False).plot(kind='bar')
plt.title('Fraud Rate by Transaction Type')
plt.xlabel('Transaction Type')
plt.ylabel('Fraud Rate (%)')
plt.xticks(rotation=45)
plt.show()

```

```
Transaction Types:  
type  
CASH_OUT    2237500  
PAYMENT     2151495  
CASH_IN     1399284  
TRANSFER     532909  
DEBIT        41432  
Name: count, dtype: int64
```



```
Fraud Rate by Transaction Type (%):  
type  
CASH_IN      0.000000  
CASH_OUT     0.183955  
DEBIT        0.000000  
PAYMENT      0.000000  
TRANSFER     0.768799  
Name: isFraud, dtype: float64
```



```

from scipy.stats import chi2_contingency

contingency_table = pd.crosstab(df['type'], df['isFraud'])
print("Contingency Table:")
print(contingency_table)

chi2, p, dof, expected = chi2_contingency(contingency_table)
print(f"\nChi-square value: {chi2}")
print(f"p-value: {p}")
print(f"Degrees of freedom: {dof}")
print(f"Statistically significant association: {p < 0.05}")

Contingency Table:
isFraud      0      1
type
CASH_IN    1399284      0
CASH_OUT   2233384  4116
DEBIT       41432      0
PAYMENT    2151495      0
TRANSFER    528812  4097

Chi-square value: 22082.53571319108
p-value: 0.0
Degrees of freedom: 4
Statistically significant association: True


numeric_df = df.select_dtypes(include=['float64', 'int64'])
plt.figure(figsize=(14, 10))

```

```

correlation_matrix = numeric_df.corr()
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm',
fmt='.2f')
plt.title('Correlation Matrix of Numeric Features')
plt.tight_layout()
plt.show()

```



```

plt.figure(figsize=(14, 6))
plt.subplot(1, 2, 1)
sns.histplot(df[df['isFraud'] == 0]['amount'], kde=True, bins=50)
plt.title('Distribution of Transaction Amounts (Non-Fraudulent)')
plt.xlabel('Amount')
plt.xlim([0, df['amount'].quantile(0.99)])
plt.subplot(1, 2, 2)
sns.histplot(df[df['isFraud'] == 1]['amount'], kde=True, bins=50,
color='red')
plt.title('Distribution of Transaction Amounts (Fraudulent)')
plt.xlabel('Amount')
plt.xlim([0, df['amount'].quantile(0.99)])

```


```

plt.tight_layout()
plt.show()


amount_stats = df.groupby('isFraud')['amount'].agg(['mean', 'median',
'min', 'max', 'std'])
print("Amount Statistics by Fraud Status:")
print(amount_stats)

plt.figure(figsize=(10, 6))
sns.boxplot(x='isFraud', y='amount', data=df, palette='Set2')
plt.title('Transaction Amount by Fraud Status')
plt.xticks([0, 1], ['Non-Fraudulent', 'Fraudulent'])
plt.ylabel('Amount')
plt.ylim([0, df['amount'].quantile(0.95)])
plt.show()

```



Amount Statistics by Fraud Status:					
	mean	median	min	max	std
isFraud					
0	1.781970e+05	74684.72	0.01	92445516.64	5.962370e+05
1	1.467967e+06	441423.44	0.00	100000000.00	2.404253e+06



```

plt.figure(figsize=(12, 6))
sns.violinplot(x='isFraud', y='amount', data=df, palette='Set2')
plt.title('Violin Plot of Transaction Amount by Fraud Status')
plt.xticks([0, 1], ['Non-Fraudulent', 'Fraudulent'])
plt.ylabel('Amount')
plt.ylim([0, df['amount'].quantile(0.95)])
plt.show()


from scipy.stats import mannwhitneyu

fraud_amounts = df[df['isFraud'] == 1]['amount']
non_fraud_amounts = df[df['isFraud'] == 0]['amount']

u_stat, p_value = mannwhitneyu(fraud_amounts, non_fraud_amounts)

print("Mann-Whitney U Test (comparing transaction amounts):")
print(f"U statistic: {u_stat}")
print(f"p-value: {p_value}")
print(f"Statistically significant difference: {p_value < 0.05}")

```




```
Mann-Whitney U Test (comparing transaction amounts):
```

```
U statistic: 41224999611.0
```

```
p-value: 0.0
```

```
Statistically significant difference: True
```

```
if 'step' in df.columns:  
    plt.figure(figsize=(15, 6))  
    fraud_over_time = df.groupby('step')['isFraud'].mean()  
    fraud_over_time.plot()  
    plt.title('Fraud Rate Over Time')  
    plt.xlabel('Time Step')  
    plt.ylabel('Fraud Rate')  
    plt.grid(True)  
    plt.show()  
  
    plt.figure(figsize=(15, 6))  
    df.groupby('step').size().plot()  
    plt.title('Transaction Volume Over Time')  
    plt.xlabel('Time Step')  
    plt.ylabel('Number of Transactions')  
    plt.grid(True)  
    plt.show()
```



```
categorical_cols =
df.select_dtypes(include=['object']).columns.tolist()
print("Categorical columns:")
for col in categorical_cols:
    print(f"{col}: {df[col].nunique()} unique values")

sample_size = 100000
if len(df) > sample_size:
    df = df.sample(n=sample_size, random_state=42)
    print(f"Sampled dataset to {len(df)} rows")

df_encoded = df.copy()
for col in categorical_cols:
    if df[col].nunique() > 50:
        encoding_map = df.groupby(col)[ 'isFraud' ].mean().to_dict()
        df_encoded[col + '_encoded'] =
df_encoded[col].map(encoding_map)
```

```

        df_encoded.drop(col, axis=1, inplace=True)
        print(f"Target encoded {col} with {df[col].nunique()}")
categories")
else:
    pass

remaining_cat_cols = [col for col in categorical_cols if col in
df_encoded.columns]
if remaining_cat_cols:
    print(f"One-hot encoding {len(remaining_cat_cols)} columns with
low cardinality")
    df_encoded = pd.get_dummies(df_encoded,
columns=remaining_cat_cols, drop_first=True)

X = df_encoded.drop('isFraud', axis=1)
y = df_encoded['isFraud']

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.3, random_state=42, stratify=y)

# Scale the features
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

print(f"Training set shape: {X_train.shape}")
print(f"Testing set shape: {X_test.shape}")

Categorical columns:
type: 5 unique values
nameOrig: 6353307 unique values
nameDest: 2722362 unique values
Sampled dataset to 100000 rows
Target encoded nameOrig with 99999 categories
Target encoded nameDest with 92914 categories
One-hot encoding 1 columns with low cardinality
Training set shape: (70000, 13)
Testing set shape: (30000, 13)

smote = SMOTE(random_state=42)
X_train_resampled, y_train_resampled =
smote.fit_resample(X_train_scaled, y_train)

print(f"Original training set shape: {X_train_scaled.shape}")
print(f"Resampled training set shape: {X_train_resampled.shape}")
print(f"Original class distribution:
{pd.Series(y_train).value_counts(normalize=True) * 100}")
print(f"Resampled class distribution:
{pd.Series(y_train_resampled).value_counts(normalize=True) * 100}")

```

```
Original training set shape: (70000, 13)
Resampled training set shape: (139802, 13)
Original class distribution: isFraud
0    99.858571
1    0.141429
Name: proportion, dtype: float64
Resampled class distribution: isFraud
0    50.0
1    50.0
Name: proportion, dtype: float64

X = df.drop(['isFraud', 'nameOrig', 'nameDest'], axis=1)
y = df['isFraud']

X = pd.get_dummies(X, columns=['type'], drop_first=True)

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

knn = KNeighborsClassifier(n_neighbors=5)
knn.fit(X_train_scaled, y_train)

y_pred_knn = knn.predict(X_test_scaled)

print("KNN Accuracy:", accuracy_score(y_test, y_pred_knn))
print("KNN Classification Report:\n",
classification_report(y_test,
y_pred_knn))

dt = DecisionTreeClassifier(
    max_depth=10,
    min_samples_split=50,
    random_state=42
)
dt.fit(X_train, y_train)

y_pred_dt = dt.predict(X_test)

print("Decision Tree Accuracy:", accuracy_score(y_test, y_pred_dt))
print("Decision Tree Classification Report:\n",
classification_report(y_test, y_pred_dt))

rf_model = RandomForestClassifier(random_state=42)
rf_model.fit(X_train_resampled, y_train_resampled)

y_pred_rf = rf_model.predict(X_test_scaled)
```

```

y_prob_rf = rf_model.predict_proba(X_test_scaled)[:, 1]

print("Random Forest Model Performance:\n")
print(f"Accuracy: {accuracy_score(y_test, y_pred_rf):.4f}")
print("\nClassification Report:")
print(classification_report(y_test, y_pred_rf))

plt.figure(figsize=(8, 6))
cm = confusion_matrix(y_test, y_pred_rf)
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.title('Confusion Matrix - Random Forest')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()


```

Random Forest Model Performance:

Accuracy: 1.0000

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	29958
1	1.00	1.00	1.00	42
accuracy			1.00	30000
macro avg	1.00	1.00	1.00	30000
weighted avg	1.00	1.00	1.00	30000



```

xgb_model = XGBClassifier(random_state=42, use_label_encoder=False,
                           eval_metric='logloss')
xgb_model.fit(X_train_resampled, y_train_resampled)

y_pred_xgb = xgb_model.predict(X_test_scaled)
y_prob_xgb = xgb_model.predict_proba(X_test_scaled)[:, 1]

print("XGBoost Model Performance:\n")
print(f"Accuracy: {accuracy_score(y_test, y_pred_xgb):.4f}")
print("\nClassification Report:")
print(classification_report(y_test, y_pred_xgb))

plt.figure(figsize=(8, 6))
cm = confusion_matrix(y_test, y_pred_xgb)
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.title('Confusion Matrix - XGBoost')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()

```


XGBoost Model Performance:

Accuracy: 1.0000

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	29958
1	1.00	1.00	1.00	42
accuracy			1.00	30000
macro avg	1.00	1.00	1.00	30000
weighted avg	1.00	1.00	1.00	30000

Confusion Matrix - XGBoost




```
feature_importance = pd.DataFrame({  
    'feature': X.columns,  
    'importance': rf_model.feature_importances_  
}).sort_values(by='importance', ascending=False)
```

```

plt.figure(figsize=(12, 8))
sns.barplot(x='importance', y='feature',
            data=feature_importance.head(15))
plt.title('Top 15 Feature Importances - Random Forest')
plt.tight_layout()
plt.show()

```



```

plt.figure(figsize=(10, 8))

fpr_rf, tpr_rf, _ = roc_curve(y_test, y_prob_rf)
auc_rf = auc(fpr_rf, tpr_rf)
plt.plot(fpr_rf, tpr_rf, label=f'Random Forest (AUC = {auc_rf:.3f})')


fpr_xgb, tpr_xgb, _ = roc_curve(y_test, y_prob_xgb)
auc_xgb = auc(fpr_xgb, tpr_xgb)
plt.plot(fpr_xgb, tpr_xgb, label=f'XGBoost (AUC = {auc_xgb:.3f})')

plt.plot([0, 1], [0, 1], 'k--')

plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve Comparison')
plt.legend(loc='lower right')

```

```
plt.grid(True)  
plt.show()
```



```
plt.figure(figsize=(10, 8))  
  
precision_rf, recall_rf, _ = precision_recall_curve(y_test, y_prob_rf)  
plt.plot(recall_rf, precision_rf, label=f'Random Forest')  
  
precision_xgb, recall_xgb, _ = precision_recall_curve(y_test,  
y_prob_xgb)  
plt.plot(recall_xgb, precision_xgb, label=f'XGBoost')  
  
plt.xlabel('Recall')  
plt.ylabel('Precision')  
plt.title('Precision-Recall Curve Comparison')  
plt.legend(loc='best')  
plt.grid(True)  
plt.show()
```

